

Linear Regression

- $y = a + bx + b_1 x_1 + b_2 x_2 + \dots$
- $y \Rightarrow$ Dependent/target(1)[1D]
- $x \Rightarrow$ independent/features(n)[2D]

```
from sklearn.linear_model import LinearRegression
import numpy as np
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error

time = np.array([5, 7, 12, 16, 20]).reshape(-1, 1)
mass = np.array([40, 120, 180, 210, 240])

mymodel = LinearRegression()
mymodel.fit(time, mass)

LinearRegression()

x = int(input("Enter the time in minute"))
result = mymodel.predict([[x]])
print("if the time is ", x, "minutes the mass is ", result[0], "grams")


Enter the time in minute 25

if the time is 25 minutes the mass is 316.7012987012987 grams

mass_model = mymodel.predict(time)
print(mass_model)

[ 72.54545455  96.96103896 158.          206.83116883 255.66233766]

import matplotlib.pyplot as plt
plt.figure(figsize=(6, 5))
plt.scatter(time, mass, label = 'Original', color='hotpink')
plt.plot(time, mass_model, label='Predicted-linear', color='k')
plt.legend()
plt.show()
```



Evaluation

R-Square

- Lower, the better

```
r2score=r2_score(time,mass_model)
print(r2score)
-816.6925282509699
```

MSE

- Lower the better

```
mse=mean_squared_error(time,mass_model)
print(mse)
25184.929870129872
```

MAE

- Lower the better

```
mae=mean_absolute_error(time,mass_model)
print(mae)
```

146.0

Linear Regression on large data

Case: Predicting the salary from age,experiance,gender,education

1. Importing libraries
2. Load data
3. Split data
4. Create and train model
5. Test the model
6. Evaluation

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import
r2_score,mean_absolute_error,mean_squared_error
from sklearn.model_selection import train_test_split

x=pd.read_csv(r"C:\Users\DELL\Downloads\Salary_EDA.csv")
x.head()
```

	Age	Gender	Education Level	Job Title	Years of Experience \
0	32.0	Male	Bachelor's	Software Engineer	5.0
1	28.0	Female	Master's	Data Analyst	3.0
2	45.0	Male	PhD	Senior Manager	15.0
3	36.0	Female	Bachelor's	Sales Associate	7.0
4	36.0	Female	Bachelor's	Sales Associate	7.0

	Salary
0	90000.0
1	65000.0
2	150000.0
3	60000.0
4	60000.0

Clean data

```
x.isnull().sum()
```

```

Age                2
Gender             4
Education Level    3
Job Title          5
Years of Experience 2
Salary            3
Gender_e           0
Education_e        0
dtype: int64

```

```

x.dropna(inplace=True)
x.isnull().sum()

```

```

Age                0
Gender             0
Education Level    0
Job Title          0
Years of Experience 0
Salary            0
Gender_e           0
Education_e        0
dtype: int64

```

Data preprocessing

```

g_e=LabelEncoder()
x['Gender_e']=g_e.fit_transform(x['Gender'])
x['Education_e']=g_e.fit_transform(x['Education Level'])
x.head()

```

	Age	Gender	Education Level	Job Title	Years of Experience
0	32.0	Male	Bachelor's	Software Engineer	5.0
1	28.0	Female	Master's	Data Analyst	3.0
2	45.0	Male	PhD	Senior Manager	15.0
3	36.0	Female	Bachelor's	Sales Associate	7.0
4	36.0	Female	Bachelor's	Sales Associate	7.0

	Salary	Gender_e	Education_e
0	90000.0	1	0
1	65000.0	0	1
2	150000.0	1	2
3	60000.0	0	0
4	60000.0	0	0

Split_ind, dep

```
X=x[['Gender_e','Education_e','Years of Experience']]
Y=x['Salary']

X_train,X_test,Y_train,Y_test=train_test_split(X,Y,train_size=0.2,random_state=42)
```

Create and Train

```
sal_model=LinearRegression()
sal_model.fit(X_train,Y_train)

LinearRegression()
```